

Computação Paralela / Computação Avançada

cap6 - 2022-11-16

Matrix Multiplication

Parallelizing with NumPy

Helmut Wolters

LIP / UC

$$C = A \cdot B = \begin{pmatrix} 38 & 92 & 92 & 52 & 15 & 76 & 66 & 45 \\ 44 & 11 & 95 & 35 & 53 & 50 & 68 & 68 \\ 65 & 54 & 44 & 37 & 32 & 42 & 20 & 30 \\ 68 & 36 & 85 & 60 & 81 & 65 & 58 & 18 \\ 95 & 71 & 55 & 31 & 23 & 15 & 64 & 40 \\ 11 & 19 & 65 & 66 & 59 & 58 & 70 & 11 \\ 55 & 29 & 32 & 71 & 96 & 59 & 54 & 87 \\ 52 & 34 & 34 & 41 & 48 & 64 & 44 & 78 \end{pmatrix} \cdot \begin{pmatrix} 55 & 17 & 94 & 40 & 74 & 63 & 63 & 89 \\ 77 & 89 & 69 & 19 & 30 & 73 & 23 & 12 \\ 40 & 59 & 53 & 29 & 28 & 57 & 29 & 22 \\ 61 & 20 & 20 & 78 & 41 & 90 & 21 & 66 \\ 46 & 45 & 79 & 12 & 85 & 76 & 89 & 23 \\ 50 & 44 & 17 & 39 & 54 & 73 & 93 & 51 \\ 59 & 57 & 94 & 65 & 27 & 79 & 81 & 31 \\ 91 & 88 & 35 & 78 & 43 & 73 & 45 & 47 \end{pmatrix}$$

$$c_{ik} = \sum_{j=1}^n a_{ij} b_{jk}$$

$$C = \begin{pmatrix} c_{11} & c_{12} & c_{13} & c_{14} & c_{15} & c_{16} & c_{17} & c_{18} \\ c_{21} & c_{22} & c_{23} & c_{24} & c_{25} & c_{26} & c_{27} & c_{28} \\ c_{31} & c_{32} & c_{33} & c_{34} & c_{35} & c_{36} & c_{37} & c_{38} \\ c_{41} & c_{42} & c_{43} & c_{44} & c_{45} & c_{46} & c_{47} & c_{48} \\ c_{51} & c_{52} & c_{53} & c_{54} & c_{55} & c_{56} & c_{57} & c_{58} \\ c_{61} & c_{62} & c_{63} & c_{64} & c_{65} & c_{66} & c_{67} & c_{68} \\ c_{71} & c_{72} & c_{73} & c_{74} & c_{75} & c_{76} & c_{77} & c_{78} \\ c_{81} & c_{82} & c_{83} & c_{84} & c_{85} & c_{86} & c_{87} & c_{88} \end{pmatrix}$$

$$C = A \cdot B = \begin{pmatrix} 38 & 92 & 92 & 52 & 15 & 76 & 66 & 45 \\ 44 & 11 & 95 & 35 & 53 & 50 & 68 & 68 \\ \color{red}{65} & \color{red}{54} & \color{red}{44} & \color{red}{37} & \color{red}{32} & \color{red}{42} & \color{red}{20} & \color{red}{33} \\ 68 & 36 & 85 & 60 & 81 & 65 & 58 & 18 \\ 95 & 71 & 55 & 31 & 23 & 15 & 64 & 40 \\ 11 & 19 & 65 & 66 & 59 & 58 & 70 & 11 \\ 55 & 29 & 32 & 71 & 96 & 59 & 54 & 87 \\ 52 & 34 & 34 & 41 & 48 & 64 & 44 & 78 \end{pmatrix} \cdot \begin{pmatrix} 55 & 17 & 94 & 40 & \color{red}{74} & 63 & 63 & 89 \\ 77 & 89 & 69 & 19 & \color{red}{30} & 73 & 23 & 12 \\ 40 & 59 & 53 & 29 & \color{red}{28} & 57 & 29 & 22 \\ 61 & 20 & 20 & 78 & \color{red}{41} & 90 & 21 & 66 \\ 46 & 45 & 79 & 12 & \color{red}{85} & 76 & 89 & 23 \\ 50 & 44 & 17 & 39 & \color{red}{54} & 73 & 93 & 51 \\ 59 & 57 & 94 & 65 & \color{red}{27} & 79 & 81 & 31 \\ 91 & 88 & 35 & 78 & \color{red}{4} & 73 & 45 & 47 \end{pmatrix}$$

$$c_{ik} = \sum_{j=1}^n a_{ij} b_{jk}$$

$$C = \begin{pmatrix} c_{11} & c_{12} & c_{13} & c_{14} & c_{15} & c_{16} & c_{17} & c_{18} \\ c_{21} & c_{22} & c_{23} & c_{24} & c_{25} & c_{26} & c_{27} & c_{28} \\ c_{31} & c_{32} & c_{33} & c_{34} & \color{red}{c_{35}} & c_{36} & c_{37} & c_{38} \\ c_{41} & c_{42} & c_{43} & c_{44} & c_{45} & c_{46} & c_{47} & c_{48} \\ c_{51} & c_{52} & c_{53} & c_{54} & c_{55} & c_{56} & c_{57} & c_{58} \\ c_{61} & c_{62} & c_{63} & c_{64} & c_{65} & c_{66} & c_{67} & c_{68} \\ c_{71} & c_{72} & c_{73} & c_{74} & c_{75} & c_{76} & c_{77} & c_{78} \\ c_{81} & c_{82} & c_{83} & c_{84} & c_{85} & c_{86} & c_{87} & c_{88} \end{pmatrix}$$

$$C = A \cdot B = \begin{pmatrix} \boxed{38} & \boxed{92} & \boxed{92} & \boxed{52} & \boxed{15} & \boxed{76} & \boxed{66} & \boxed{45} \\ \boxed{44} & \boxed{11} & \boxed{95} & \boxed{35} & \boxed{53} & \boxed{50} & \boxed{68} & \boxed{68} \\ \boxed{65} & \boxed{54} & \boxed{44} & \boxed{37} & \boxed{32} & \boxed{42} & \boxed{20} & \boxed{30} \\ \boxed{68} & \boxed{36} & \boxed{85} & \boxed{60} & \boxed{81} & \boxed{65} & \boxed{58} & \boxed{18} \\ \boxed{95} & \boxed{71} & \boxed{55} & \boxed{31} & \boxed{23} & \boxed{15} & \boxed{64} & \boxed{40} \\ \boxed{11} & \boxed{19} & \boxed{65} & \boxed{66} & \boxed{59} & \boxed{58} & \boxed{70} & \boxed{11} \\ \boxed{55} & \boxed{29} & \boxed{32} & \boxed{71} & \boxed{96} & \boxed{59} & \boxed{54} & \boxed{87} \\ \boxed{52} & \boxed{34} & \boxed{34} & \boxed{41} & \boxed{48} & \boxed{64} & \boxed{44} & \boxed{78} \end{pmatrix} \cdot \begin{pmatrix} 55 & 17 & 94 & 40 & 74 & 63 & 63 & 89 \\ 77 & 89 & 69 & 19 & 30 & 73 & 23 & 12 \\ 40 & 59 & 53 & 29 & 28 & 57 & 29 & 22 \\ 61 & 20 & 20 & 78 & 41 & 90 & 21 & 66 \\ 46 & 45 & 79 & 12 & 85 & 76 & 89 & 23 \\ 50 & 44 & 17 & 39 & 54 & 73 & 93 & 51 \\ 59 & 57 & 94 & 65 & 27 & 79 & 81 & 31 \\ 91 & 88 & 35 & 78 & 43 & 73 & 45 & 47 \end{pmatrix}$$

$$c_{ik} = \sum_{j=1}^n a_{ij} b_{jk}$$

$$C = \begin{pmatrix} \boxed{c_{11}} & \boxed{c_{12}} & \boxed{c_{13}} & \boxed{c_{14}} & \boxed{c_{15}} & \boxed{c_{16}} & \boxed{c_{17}} & \boxed{c_{18}} \\ \boxed{c_{21}} & \boxed{c_{22}} & \boxed{c_{23}} & \boxed{c_{24}} & \boxed{c_{25}} & \boxed{c_{26}} & \boxed{c_{27}} & \boxed{c_{28}} \\ \boxed{c_{31}} & \boxed{c_{32}} & \boxed{c_{33}} & \boxed{c_{34}} & \boxed{c_{35}} & \boxed{c_{36}} & \boxed{c_{37}} & \boxed{c_{38}} \\ \boxed{c_{41}} & \boxed{c_{42}} & \boxed{c_{43}} & \boxed{c_{44}} & \boxed{c_{45}} & \boxed{c_{46}} & \boxed{c_{47}} & \boxed{c_{48}} \\ \boxed{c_{51}} & \boxed{c_{52}} & \boxed{c_{53}} & \boxed{c_{54}} & \boxed{c_{55}} & \boxed{c_{56}} & \boxed{c_{57}} & \boxed{c_{58}} \\ \boxed{c_{61}} & \boxed{c_{62}} & \boxed{c_{63}} & \boxed{c_{64}} & \boxed{c_{65}} & \boxed{c_{66}} & \boxed{c_{67}} & \boxed{c_{68}} \\ \boxed{c_{71}} & \boxed{c_{72}} & \boxed{c_{73}} & \boxed{c_{74}} & \boxed{c_{75}} & \boxed{c_{76}} & \boxed{c_{77}} & \boxed{c_{78}} \\ \boxed{c_{81}} & \boxed{c_{82}} & \boxed{c_{83}} & \boxed{c_{84}} & \boxed{c_{85}} & \boxed{c_{86}} & \boxed{c_{87}} & \boxed{c_{88}} \end{pmatrix}$$

Example:

- n lines and n rows of floating numbers with $n=1024$
- Floating number occupies 8 bytes
- Matrix: $1024 \times 1024 \times 8 = 8$ MegaBytes
- $np = 128$ processes
- Master has to send $1/128$ of A and the whole matrix B to each process
 - A : $127 \times \text{Send} \approx 8$ MegaBytes over the network
1 **Scatter** ≈ 8 MegaBytes over the network
(no change, but much faster due to optimization by MPI)
 - B : $127 \times \text{Send} = 127 \times 8$ Megabytes ≈ 1 GigaByte over the network (!)
1 **Broadcast**: only ≈ 8 MegaBytes over the network !
- Master receives $1/128$ of C from each process
 - $127 \text{ Recv} \approx 8$ MegaBytes over the network
 - 1 **Gather** ≈ 8 MegaBytes over the network
(no change, but **very** much faster due to optimization by MPI)

$$C = A \cdot B = \begin{pmatrix} \boxed{38} & \boxed{92} & \boxed{92} & \boxed{52} & \boxed{15} & \boxed{76} & \boxed{66} & \boxed{45} \\ \boxed{44} & \boxed{11} & \boxed{95} & \boxed{35} & \boxed{53} & \boxed{50} & \boxed{68} & \boxed{68} \\ \boxed{65} & \boxed{54} & \boxed{44} & \boxed{37} & \boxed{32} & \boxed{42} & \boxed{20} & \boxed{30} \\ \boxed{68} & \boxed{36} & \boxed{85} & \boxed{60} & \boxed{81} & \boxed{65} & \boxed{58} & \boxed{18} \\ \boxed{95} & \boxed{71} & \boxed{55} & \boxed{31} & \boxed{23} & \boxed{15} & \boxed{64} & \boxed{40} \\ \boxed{11} & \boxed{19} & \boxed{65} & \boxed{66} & \boxed{59} & \boxed{58} & \boxed{70} & \boxed{11} \\ \boxed{55} & \boxed{29} & \boxed{32} & \boxed{71} & \boxed{96} & \boxed{59} & \boxed{54} & \boxed{87} \\ \boxed{52} & \boxed{34} & \boxed{34} & \boxed{41} & \boxed{48} & \boxed{64} & \boxed{44} & \boxed{78} \end{pmatrix} \cdot \begin{pmatrix} 55 & 17 & 94 & 40 & 74 & 63 & 63 & 89 \\ 77 & 89 & 69 & 19 & 30 & 73 & 23 & 12 \\ 40 & 59 & 53 & 29 & 28 & 57 & 29 & 22 \\ 61 & 20 & 20 & 78 & 41 & 90 & 21 & 66 \\ 46 & 45 & 79 & 12 & 85 & 76 & 89 & 23 \\ 50 & 44 & 17 & 39 & 54 & 73 & 93 & 51 \\ 59 & 57 & 94 & 65 & 27 & 79 & 81 & 31 \\ 91 & 88 & 35 & 78 & 43 & 73 & 45 & 47 \end{pmatrix}$$

Scatter Broadcast

$$c_{ik} = \sum_{j=1}^n a_{ij} b_{jk}$$

Gather

$$C = \begin{pmatrix} \boxed{c_{11}} & \boxed{c_{12}} & \boxed{c_{13}} & \boxed{c_{14}} & \boxed{c_{15}} & \boxed{c_{16}} & \boxed{c_{17}} & \boxed{c_{18}} \\ \boxed{c_{21}} & \boxed{c_{22}} & \boxed{c_{23}} & \boxed{c_{24}} & \boxed{c_{25}} & \boxed{c_{26}} & \boxed{c_{27}} & \boxed{c_{28}} \\ \boxed{c_{31}} & \boxed{c_{32}} & \boxed{c_{33}} & \boxed{c_{34}} & \boxed{c_{35}} & \boxed{c_{36}} & \boxed{c_{37}} & \boxed{c_{38}} \\ \boxed{c_{41}} & \boxed{c_{42}} & \boxed{c_{43}} & \boxed{c_{44}} & \boxed{c_{45}} & \boxed{c_{46}} & \boxed{c_{47}} & \boxed{c_{48}} \\ \boxed{c_{51}} & \boxed{c_{52}} & \boxed{c_{53}} & \boxed{c_{54}} & \boxed{c_{55}} & \boxed{c_{56}} & \boxed{c_{57}} & \boxed{c_{58}} \\ \boxed{c_{61}} & \boxed{c_{62}} & \boxed{c_{63}} & \boxed{c_{64}} & \boxed{c_{65}} & \boxed{c_{66}} & \boxed{c_{67}} & \boxed{c_{68}} \\ \boxed{c_{71}} & \boxed{c_{72}} & \boxed{c_{73}} & \boxed{c_{74}} & \boxed{c_{75}} & \boxed{c_{76}} & \boxed{c_{77}} & \boxed{c_{78}} \\ \boxed{c_{81}} & \boxed{c_{82}} & \boxed{c_{83}} & \boxed{c_{84}} & \boxed{c_{85}} & \boxed{c_{86}} & \boxed{c_{87}} & \boxed{c_{88}} \end{pmatrix}$$

$$C = A \cdot B = \begin{pmatrix} \boxed{\begin{matrix} 38 & 92 & 92 & 52 & 15 & 76 & 66 & 45 \\ 44 & 11 & 95 & 35 & 53 & 50 & 68 & 68 \\ 65 & 54 & 44 & 37 & 32 & 42 & 20 & 30 \\ 68 & 36 & 85 & 60 & 81 & 65 & 58 & 18 \end{matrix}} \\ \boxed{\begin{matrix} 95 & 71 & 55 & 31 & 23 & 15 & 64 & 40 \\ 11 & 19 & 65 & 66 & 59 & 58 & 70 & 11 \\ 55 & 29 & 32 & 71 & 96 & 59 & 54 & 87 \\ 52 & 34 & 34 & 41 & 48 & 64 & 44 & 78 \end{matrix}} \end{pmatrix} \cdot \begin{pmatrix} \boxed{\begin{matrix} 55 & 17 & 94 & 40 & 74 & 63 & 63 & 89 \\ 77 & 89 & 69 & 19 & 30 & 73 & 23 & 12 \\ 40 & 59 & 53 & 29 & 28 & 57 & 29 & 22 \\ 61 & 20 & 20 & 78 & 41 & 90 & 21 & 66 \\ 46 & 45 & 79 & 12 & 85 & 76 & 89 & 23 \\ 50 & 44 & 17 & 39 & 54 & 73 & 93 & 51 \\ 59 & 57 & 94 & 65 & 27 & 79 & 81 & 31 \\ 91 & 88 & 35 & 78 & 43 & 73 & 45 & 47 \end{matrix}} \end{pmatrix}$$

Send/Recv Send/Recv

$$c_{ik} = \sum_{j=1}^n a_{ij} b_{jk}$$

Gather

$$C = \begin{pmatrix} \boxed{\begin{matrix} c_{11} & c_{12} & c_{13} & c_{14} \\ c_{21} & c_{22} & c_{23} & c_{24} \\ c_{31} & c_{32} & c_{33} & c_{34} \\ c_{41} & c_{42} & c_{43} & c_{44} \end{matrix}} \boxed{\begin{matrix} c_{15} & c_{16} & c_{17} & c_{18} \\ c_{25} & c_{26} & c_{27} & c_{28} \\ c_{35} & c_{36} & c_{37} & c_{38} \\ c_{45} & c_{46} & c_{47} & c_{48} \end{matrix}} \\ \boxed{\begin{matrix} c_{51} & c_{52} & c_{53} & c_{54} \\ c_{61} & c_{62} & c_{63} & c_{64} \\ c_{71} & c_{72} & c_{73} & c_{74} \\ c_{81} & c_{82} & c_{83} & c_{84} \end{matrix}} \boxed{\begin{matrix} c_{55} & c_{56} & c_{57} & c_{58} \\ c_{65} & c_{66} & c_{67} & c_{68} \\ c_{75} & c_{76} & c_{77} & c_{78} \\ c_{85} & c_{86} & c_{87} & c_{88} \end{matrix}} \end{pmatrix}$$

Example:

- $n=1024$, Matrix: $1024 \times 1024 \times 8 = 8$ MegaBytes, $np = 128$ processes
- Master has to send $1/64$ of A and $1/64$ of B to each process
 - A: $127 \times \text{Send} \approx 8$ MegaBytes over the network
 - B: $127 \times \text{Send} \approx 8$ MegaBytes over the network
- Master receives $1/128$ of C from each process
 - $127 \text{ Recv} \approx 8$ MegaBytes over the network
 - $1 \text{ Gather} \approx 8$ MegaBytes over the network
(no change, but much faster due to optimization by MPI)
- **And more:** each process needs much less memory:
 - Instead of $1/128$ of A and whole B: $1/64$ of A and $1/64$ of B
 - Instead of 64 kiloBytes + 8 MegaBytes: $2 \times 128 \text{ kiloBytes} = 256 \text{ kiloBytes}$
- Doesn't sound much,
 - but now go from $np=128$ processes to 1024 processes
 - and from $n=1024$ to $n = 1024 \times 1024 \approx 10^6$
 - **This grows exponentially !**

Scatter

Scatter

$$c_{ik} = \sum_{j=1}^n a_{ij}b_{jk}$$

Gather

$$C = \begin{pmatrix} c_{11} & c_{12} & c_{13} & c_{14} & c_{15} & c_{16} & c_{17} & c_{18} \\ c_{21} & c_{22} & c_{23} & c_{24} & c_{25} & c_{26} & c_{27} & c_{28} \\ c_{31} & c_{32} & c_{33} & c_{34} & c_{35} & c_{36} & c_{37} & c_{38} \\ c_{41} & c_{42} & c_{43} & c_{44} & c_{45} & c_{46} & c_{47} & c_{48} \\ c_{51} & c_{52} & c_{53} & c_{54} & c_{55} & c_{56} & c_{57} & c_{58} \\ c_{61} & c_{62} & c_{63} & c_{64} & c_{65} & c_{66} & c_{67} & c_{68} \\ c_{71} & c_{72} & c_{73} & c_{74} & c_{75} & c_{76} & c_{77} & c_{78} \\ c_{81} & c_{82} & c_{83} & c_{84} & c_{85} & c_{86} & c_{87} & c_{88} \end{pmatrix}$$

$$\begin{pmatrix} 55 & 17 & 94 & 40 & 74 & 63 & 63 & 89 \\ 77 & 89 & 69 & 19 & 30 & 73 & 23 & 12 \\ 40 & 59 & 53 & 29 & 28 & 57 & 29 & 22 \\ 61 & 20 & 20 & 78 & 41 & 90 & 21 & 66 \\ 46 & 45 & 79 & 12 & 85 & 76 & 89 & 23 \\ 50 & 44 & 17 & 39 & 54 & 73 & 93 & 51 \\ 59 & 57 & 94 & 65 & 27 & 79 & 81 & 31 \\ 91 & 88 & 35 & 78 & 43 & 73 & 45 & 47 \end{pmatrix}$$

Scatter

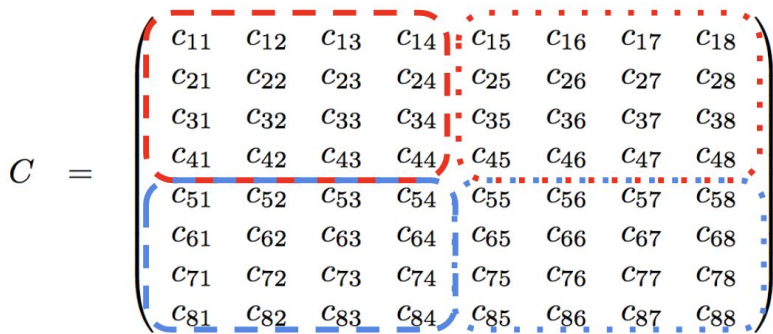
This is not as trivial as it looks like!

The Scatter method does not send a matrix,
it just sends a stream of bytes,
following the numbers line by line:

55 17 94 40 74 63 63 89 77 89 69 19 30 73 23 12 40 59 53 29 28 57 29 22 61 20 20 78 41 90 21 66 46 45 79 12 85 76 ...

- We could use “scatter” (lowercase s), but this destroys all the efficiency when coming to large matrices
- Solution:
 Transpose matrix B
- This is a costly procedure when B is large, so
 in production environment, you would store B transposed from the beginning
- But in our exercises, we will just transpose it, so we are sure we know what we are doing...

Gather



This is even worse!

The Gather method does not send/receive pieces of a matrix, every process just sends a stream of bytes for each chunk, and the master receives all the chunks in the sequence of the processes:

$c_{11} \ c_{12} \ c_{13} \ c_{14} \ c_{21} \ c_{22} \ c_{23} \ c_{24} \ c_{31} \ c_{32} \ c_{33} \ c_{34} \ c_{41} \ c_{42} \ c_{43} \ c_{44} \ c_{15} \ c_{16} \ c_{17} \ c_{18} \ c_{25} \ c_{26} \ c_{27} \ c_{28} \ c_{35} \ c_{36} \ c_{37} \ c_{38} \ c_{45} \ c_{46} \ c_{47} \ \dots$

- Solution: numpy is your friend - it has all the methods you need for this type of problems
- Concrete solution for reorganizing the chunks from the incoming linear stream:
`np.block`

Numpy methods

```
x = np.linspace(1,n**2,n**2)
```

```
A = x.reshape((n,n))
```

```
B = np.linspace(10,10*n**2,n**2).reshape((n,n))
```

```
C = np.matmul(A, B)
```

Strategy:

- Start with $n=8$, single core, just a straight python program with numpy
- Create two matrices A and B, and determine $C = \text{np.matmul}(A, B)$ as reference
- Transform both A and B in linear arrays with reshape and play with np.reshape and np.block to understand the details of how it works:
 - Create the chunks you will need for the parallel process
 - Do the matmul on each chunk
 - Join the result back into one data stream
 - Transform the stream into the correct resulting matrix
- Develop your parallel code with the exact same example and make it work
- That is: run it, check it, repair it, until you got the chunks right
- Make it work for any n for which n^2 is a multiple of np.

MPI methods for exchanging numpy arrays in messages:

- Same names, but starting with a capital letter:
 - `comm.Send (numpy_array, dest=p)`
 - `comm.Recv (numpy_array, source=0)`
 - etc
- Note that the `numpy_array`, on the receiving side,
 - **Must exist and have been created with the necessary space to hold the array it will receive**
 - Typically with something like: `numpy_array = np.empty (numData, dtype=float)`
- Have a look at the example programs called *numpy?.py* for detailed examples of all the methods
- These methods send a bit stream of the object that then needs to be interpreted correctly on the receiving end
- The advantage is that this is much faster than sending a generic python object that contains a lot of overhead of metadata, and for array calculations we usually have just big sequences of floating point numbers that can easily be reordered through numpy methods.
- Have a look at the example programs called *matrix?.py*

Numpy methods

- a good starting point for an introduction to the *Numpy* package is the tutorial at [NumPy quickstart — NumPy v1.24.dev0 Manual](#)