

Parallel Computing

2022/2023 1st semester

Exercise Sheet #6 2022-11-16

Create a parallel processing code that performs the multiplication of two matrices A and B of dimension $n \times n$. Read n from the command line and use a random process to create the components of the matrices.

Note that, while testing the code, it is recommended to create matrices with systematic values for their components instead of a random process, in order to simplify debugging.

1. Distribute the lines of matrix A equally over $n_p = 4$ participating processes. This way, the master process has to send the corresponding part of matrix A and the whole matrix B to each of the slave processes, and the slaves send back the corresponding lines of the result matrix.
2. Distribute equal parts of both matrix A and matrix B over the processes. If the number of processes n_p involved is a square number so that $n_p = m^2$, you can divide the lines of matrix A and the columns of matrix B in m equally sized segments and thus have each slave process calculate one segment of the result with $\frac{n}{m}$ lines and $\frac{n}{m}$ columns. Make sure that each slave process receives only the parts of A and B that it needs for its part of the calculation. This reduces drastically the amount of data transferred through MPI.

Note that for an efficient transmission of subsequent columns of matrix B , it might be useful to use the transposed matrix (`np.transpose`).

In order for the master to recombine the resulting matrix from the segments sent by the slaves, you might want to have a look at the block method of numpy – `help(np.block)` .

Begin with $n_p = n = 4$, and then make the code work for numbers of processes n_p that are powers of 4 (4, 16, 64, 256, ...) and matrices with a dimension n such that n^2 is a multiple of n_p .

Submit the codes of your programs (1 to 2) at <https://trixi.coimbra.lip.pt/cap> until the beginning of the next lectures.